

Task 3.6 Summarising & Cleaning Data in SQL

By Lee Heng Chuah

In this task you will calculate some descriptive statistics using the MIN, MAX, AVG, COUNT, SUM and MODE() aggregates discussed in this Exercise, and you will reflect on what you learned about data profiling back in Exercise 1.5: Data Profiling & Integrity.

Directions:

Rockbuster's database engineers have loaded some new data into the database, and your manager has asked you to clean and profile it. Follow the instructions below to complete their request:

1. **Check for and clean dirty data:** Find out if the film table and the customer table contain any dirty data, specifically non-uniform or duplicate data, or missing values. Create a new "Answer 3.6" document and copy-paste your queries into it. Next to each query write 2 to 3 sentences explaining how you would clean the data (even if the data is not dirty)

Find the duplicate data (film):

Dashboard Properties SQL Statistics Dependencies Dependents Rockbuster/postgres@PostgreSQL 14*

Rockbuster/postgres@PostgreSQL 14

No limit

Query Query History Scratch Pad x

```

1 -- Only show those record that duplicate data (based on columns selected)
2
3 SELECT title,
4       release_year,
5       language_id,
6       rental_duration,
7       COUNT(*)
8 FROM film
9 GROUP BY title,
10        release_year,
11        language_id,
12        rental_duration
13 HAVING COUNT(*) > 1; -- no result set means we have no duplicates

```

Data output Messages Explain x Notifications

	title	release_year	language_id	rental_duration	count
	character varying (255)	integer	smallint	smallint	bigint

Total rows: 0 of 0 Query complete 00:00:00.083 Ln 13, Col 66

Find the duplicate data (customer):

The screenshot shows a PostgreSQL query editor interface. The top navigation bar includes 'Dashboard', 'Properties', 'SQL', 'Statistics', 'Dependencies', and 'Dependents'. The current connection is 'Rockbuster/postgres@PostgreSQL 14*'. The query editor has a toolbar with icons for file operations, query execution, and settings. The 'Query' tab is active, showing the following SQL query:

```
1  -- Only show those record that duplicate data (based on column selected)
2
3  SELECT store_id,
4         first_name,
5         last_name,
6         email,
7         address_id,
8         active,
9         COUNT(*)
10 FROM customer
11 GROUP BY store_id,
12          first_name,
13          last_name,
14          email,
15          address_id,
16          active
17 HAVING COUNT (*) > 1; -- no result set means we have no duplicates
18
```

The 'Data output' tab is also visible, showing the column definitions for the query result:

store_id	first_name	last_name	email	active	count
smallint	character varying (45)	character varying (45)	character varying (50)	smallint	bigint

The status bar at the bottom indicates 'Total rows: 0 of 0', 'Query complete 00:00:00.050', and 'Ln 18, Col 1'.

There are no duplications on both film and customer data, as we can see from the Data Output above.

As described in Task 3.6, there are two ways that we can fix them if we find duplicate records in our database:

- Create a virtual table, known as a “view”, where you select only unique records – standard way
- Delete the duplicate record from the table or view

2. Summarise your data: Use SQL to calculate descriptive statistics for both the film table and the customer table. For numerical columns, this means finding the minimum, maximum and average values. For non-

numerical columns, calculate the mode value. Copy-paste your SQL queries and their outputs into your answers document.

Film – numerical columns

Dashboard Properties SQL Statistics Dependencies Dependents [Rockbuster/postgres@PostgreSQL 14*](#)

Rockbuster/postgres@PostgreSQL 14

Query Query History

```

1 SELECT MIN(film_id) AS min_film_id,
2        MAX(film_id) AS max_film_id,
3        AVG(film_id) AS avg_film_id,
4        COUNT(film_id) AS count_film_id,
5        MIN(release_year) AS min_release_year,
6        MAX(release_year) AS max_release_year,
7        AVG(release_year) AS avg_release_year,
8        COUNT(release_year) AS count_release_year,
9        MIN(language_id) AS min_language_id,
10       MAX(language_id) AS max_language_id,
11       AVG(language_id) AS avg_language_id,
12       COUNT(language_id) AS count_language_id,
13       MIN(rental_duration) AS min_rental_duration,
14       MAX(rental_duration) AS max_rental_duration,
15       AVG(rental_duration) AS avg_rental_duration,
16       COUNT(rental_duration) AS count_rental_duration,
17       MIN(rental_rate) AS min_rental_rate,
18       MAX(rental_rate) AS max_rental_rate,
19       AVG(rental_rate) AS avg_rental_rate,
20       COUNT(rental_rate) AS count_rental_rate,
21       MIN(length) AS min_length,
22       MAX(length) AS max_length,
23       AVG(length) AS avg_length,
24       COUNT(length) AS count_length,
25       MIN(replacement_cost) AS min_replacement_cost,
26       MAX(replacement_cost) AS max_replacement_cost,
27       AVG(replacement_cost) AS avg_replacement_cost,
28       COUNT(replacement_cost) AS count_replacement_cost,
29       COUNT(*) AS count_rows
30 FROM film
  
```

Film – numerical columns – Output (Note: The output has been split into a few screens due to the length of information)

min_film_id integer	max_film_id integer	avg_film_id numeric	count_film_id bigint
1	1000	500.5	1000

min_release_year integer	max_release_year integer	avg_release_year numeric	count_release_year bigint
2006	2006	2006	1000

min_language_id smallint	max_language_id smallint	avg_language_id numeric	count_language_id bigint
1	1	1	1000

min_rental_duration smallint	max_rental_duration smallint	avg_rental_duration numeric	count_rental_duration bigint
3	7	4.985	1000

min_rental_rate numeric	max_rental_rate numeric	avg_rental_rate numeric	count_rental_rate bigint
0.99	4.99	2.98	1000

min_replacement_cost numeric	max_replacement_cost numeric	avg_replacement_cost numeric	count_replacement_cost bigint	count_rows bigint
9.99	29.99	19.984	1000	1000

min_length	max_length	avg_length	count_length
smallint	smallint	numeric	bigint
46	185	115.272	1000

Customer – numerical columns

DashboardPropertiesSQLStatisticsDependenciesDependentsRockbuster/postgres@PostgreSQL 14*

Rockbuster/postgres@PostgreSQL 14

Query

Query History

Scratch Pad

1SELECT MIN(customer_id) AS min_customer_id,

2MAX(customer_id) AS max_customer_id,

3AVG(customer_id) AS avg_customer_id,

4COUNT(customer_id) AS count_customer_id,

5MIN(store_id) AS min_store_id,

6MAX(store_id) AS max_store_id,

7AVG(store_id) AS avg_store_id,

8COUNT(store_id) AS count_store_id,

9MIN(address_id) AS min_address_id,

10MAX(address_id) AS max_address_id,

11AVG(address_id) AS avg_address_id,

12COUNT(address_id) AS count_address_id,

13COUNT(*) AS count_rows

14FROM customer

Data outputMessagesExplainNotifications

min_customer_idinteger1599

max_customer_idinteger300

avg_customer_idnumeric599

count_customer_idbigint1

min_store_idsmallint21.45575959933

max_store_idsmallint599

avg_store_idnumeric605

count_store_idbigint5

min_address_idsmallint304.72454090150

max_address_idsmallint599

avg_address_idnumeric599

count_address_idbigint599

count_rowsbigint599

Total rows: 1 of 1Query complete 00:00:00.153Ln 13, Col 30

Film – non-numerical columns

Dashboard Properties SQL Statistics Dependencies Dependents Rockbuster/postgres@PostgreSQL 14*

Rockbuster/postgres@PostgreSQL 14

No limit

Query Query History Scratch Pad

```

1 SELECT mode() WITHIN GROUP (ORDER BY title)
2     AS mode_title_value,
3     mode() WITHIN GROUP (ORDER BY rating)
4     AS mode_rating_value,
5     mode() WITHIN GROUP (ORDER BY special_features)
6     AS mode_special_feature_value
7 FROM film;
8

```

Data output Messages Explain × Notifications

	mode_title_value <small>character varying</small>	mode_rating_value <small>mpaa_rating</small>	mode_special_feature_value <small>text[]</small>
1	Academy Dinosaur	PG-13	{Trailers,Commentaries,'Behind the Scenes'}

Total rows: 1 of 1 Query complete 00:00:00.079 Ln 8, Col 1

Customer – non-numerical columns

Dashboard Properties SQL Statistics Dependencies Dependents **Rockbuster/postgres@PostgreSQL 14*** + x

Rockbuster/postgres@PostgreSQL 14

Query Query History Scratch Pad x

```

1 SELECT mode() WITHIN GROUP (ORDER BY first_name)
2     AS mode_first_name_value,
3     mode() WITHIN GROUP (ORDER BY last_name)
4     AS mode_last_name_value,
5     mode() WITHIN GROUP (ORDER BY email)
6     AS mode_email_value,
7     mode() WITHIN GROUP (ORDER BY activebool)
8     AS mode_activebool_value
9 FROM customer;
10

```

Data output Messages Explain x Notifications

	mode_first_name_value character varying	mode_last_name_value character varying	mode_email_value character varying	mode_activebool_value boolean
1	Jamie	Abney	aaron.selby@sakilacustomer.org	true

Total rows: 1 of 1 Query complete 00:00:00.117 Ln 6, Col 25

- Reflect on your work:** Back in Achievement 1 you learned about data profiling in Excel. Based on your previous experience, which tool (Excel or SQL) do you think is more effective for data profiling, and why? Consider their respective functions, ease of use, and speed. Write a short paragraph in the running document that you have started.

Back in Achievement 1, it took me approximate 2 weeks to find the duplication and missing data with Excel. In this achievement, it took me less than 3 hours to find the repetition and missing data using the SQL function. Excel took me a lengthy time to average duplicated data and eliminate the duplication of data, while SQL could average the data by using the function. Clearly, the SQL is faster, easier and safer than Excel.

- Save your “Answer 3.6” document as a PDF and upload it here for your tutor to review.